

# Accelerated Chi-Squared

May 3, 2019

## 0.1 Project Description

In this project, we aim to calculate the following equation on multiple GPUs using PyTorch and find its lowest value on a grid where  $i, j$  go up to 80, 60 respectively. Ultimately, we wish to compare SOM mapping time on a GPU and a CPU.

$$\chi_{i,j}^2 = \sum_{k=1}^{11} (G^k - W_{i,j}^k)$$

Therefore, for purposes of programming with PyTorch,  $W$  is a single tensor with depth, height, width of 11, 60, 80 and  $G$ 's are a set of tensors with depth, height, width of 11, 1, 1.

Author: Miladiouss (Milad Pourrahmani)

Email : mpourrah@uci.edu / miladiouss@gmail.com

Date : May 3, 2019

## 0.2 Imports and GPU Test

```
In [2]: #-----
# IPython
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
from IPython import display
#-----

from time import time
import numpy as np
import math
from numba import vectorize, cuda

In [3]: from datetime import datetime
timeFormat = '%a %b %d %Y at %H:%M:%S'
import socket
deviceName = socket.gethostname()
import getpass
userName = getpass.getuser()

t1 = datetime.now()
print(t1.strftime(timeFormat) + ' on ' + deviceName + ' by ' + userName)
```

```

# GPU Detection Tests
from torch.cuda import device_count
import torch
import gpustat

if False:
    print("""
    If pytorch raises "RuntimeError: cuda runtime error (30)" after suspension,
    run the following commands in linux terminal:

    sudo rmmod nvidia_uvm
    sudo rmmod nvidia
    sudo modprobe nvidia
    sudo modprobe nvidia_uvm

    Then restart Jupyter kernel
    """)

print('gpustat output:\n')
gpustat.print_gpustat()
print('\n\n')

numGPUs = torch.cuda.device_count()
print('=== Name of {} GPUs by PyTorch ===\n'.format(numGPUs))
for i in range(numGPUs):
    print('[{}]: {}'.format(i, torch.cuda.get_device_name(i)))
print('=====')

torch.cuda.empty_cache()

```

Fri May 03 2019 at 02:06:06 on neuron by miladiouss  
gpustat output:

```

neuron Fri May 3 02:06:06 2019
[0] GeForce 940MX      | 47'C,  0 % |      0 / 2004 MB |
[1] TITAN V           | 37'C, 10 % |      0 / 12036 MB |

```

=== Name of 2 GPUs by PyTorch ===

```

[0]: TITAN V
[1]: GeForce 940MX
=====

```

### 0.3 PyTorch Implementation and Timing

```
In [90]: for N in (10, 100, 1000, 10000, 100000, 1000000):
        _ = torch.manual_seed(42)
        vs = 11
        #     N = 100 # Suggestion for CPU/GPU: 1e4/1e6
        bs = 1000 # Suggestion for GPU: 1000
        printMod = ['t_tot', 'all'][1]
        #=====

        # Imports
        from time import time

        import torch
        import torch.nn as nn
        import torch.nn.functional as F
        from torch.utils.data import Dataset, DataLoader

        # Device managements
        deviceID = 0
        device = torch.device('cuda:{}'.format(deviceID) if torch.cuda.is_available() else
        if False:
            device = 'cpu'
            torch.set_num_threads = 1
        else:
            if printMod == 'all':
                print('device name      : {}'.format(torch.cuda.get_device_name(device)
        if printMod == 'all':
            print('device              : {}'.format(device))

        # Dataset
        class RandomDataset(Dataset):

            def __init__(self, size):
                self.len = size

            def __getitem__(self, index):
                return torch.randn((vs, 1, 1))

            def __len__(self):
                return self.len

        valset = RandomDataset(N)
        valLoader = DataLoader(valset, batch_size=bs, pin_memory=False)

        # Main model
        class Model(nn.Module):
            def __init__(self, W_depth = vs, W_height = 80, W_width = 60, feed = None, bs
```

```

        super(Model, self).__init__()
        self.W = torch.rand((W_depth, W_height, W_width))
        if feed is not None:
            self.W = feed
        self.W = self.W.to(device)
        self.localGs = torch.randn((bs, vs, 1, 1))
        self.localGs = self.localGs.to(device)

    def chiFunc(self, Gs):
        '''Works on batches'''
        output = torch.add(self.W, -Gs)**2
        output = output.sum(1)
        output = torch.min(output.flatten(1), 1)
        # Returns lowest values and corresponding locations
        return output

    def forward(self, Gs, local = False):
        if Gs is None:
            Gs = self.localGs
        return self.chiFunc(Gs)

# Place model on device
model = Model()
if True and device != 'cpu':
    model = torch.nn.DataParallel(model, device_ids=[deviceID])
else:
    model = model.to(device)

# Timing setup
t_calc = 0
t_trans = 0
t_tot = 0

t1_tot = time()

# Main calculations

# Mock data is created on GPU
if True:
    if printMod == 'all':
        print("Data Management : Generated on device")
    for i in range(N//bs):
        t1_calc = time()
        output = model(None)
        t2_calc = time()
        t_calc += t2_calc - t1_calc

```

```

# Mock data is transfered to GPU
else:
    if printMod == 'all':
        print("Data Management      : Transferred to device")
        for i, Gbatch in enumerate(valLoader):
            t1_trans = time()
            Gbatch = Gbatch.to(device)
            t2_trans = time()
            t_trans += t2_trans - t1_trans

            t1_calc = time()
            output = model(Gbatch)
            t2_calc = time()
            t_calc += t2_calc - t1_calc

        t2_tot = time()
        t_tot = t2_tot - t1_tot

# Reports
if printMod == 'all':
    print('Load Time          : {: e} s per {:.0e}'.format(t_tot - t_trans - t_calc, N))
    print('Transfer Time       : {: e} s per {:.0e}'.format(t_trans, N))
    print('Calculation Time      : {: e} s per {:.0e}'.format(t_calc, N))
    print('Total          Time    : {: e} s per {:.0e}'.format(t_tot, N))
    print("=====")
if printMod == 't_tot':
    print(t_tot, end=', ')

```

```

device name      : TITAN V
device           : cuda:0
Data Management  : Generated on device
Load Time        : 5.674362e-05 s per 1e+01
Transfer Time     : 0.000000e+00 s per 1e+01
Calculation Time  : 0.000000e+00 s per 1e+01
Total           : 5.674362e-05 s per 1e+01
Shooby's CPU Time : 9.000000e+00 s per 1e+04

```

```

=====
device name      : TITAN V
device           : cuda:0
Data Management  : Generated on device
Load Time        : 6.365776e-05 s per 1e+02
Transfer Time     : 0.000000e+00 s per 1e+02
Calculation Time  : 0.000000e+00 s per 1e+02
Total           : 6.365776e-05 s per 1e+02
Shooby's CPU Time : 9.000000e+00 s per 1e+04

```

```

=====
device name      : TITAN V

```

```

device           : cuda:0
Data Management  : Generated on device
Load Time        : 9.822845e-05 s per 1e+03
Transfer Time     : 0.000000e+00 s per 1e+03
Calculation Time  : 4.632473e-04 s per 1e+03
Total Time       : 5.614758e-04 s per 1e+03
Shooby's CPU Time : 9.000000e+00 s per 1e+04

```

```

=====
device name      : TITAN V
device           : cuda:0
Data Management  : Generated on device
Load Time        : 7.224083e-05 s per 1e+04
Transfer Time     : 0.000000e+00 s per 1e+04
Calculation Time  : 1.677513e-03 s per 1e+04
Total Time       : 1.749754e-03 s per 1e+04
Shooby's CPU Time : 9.000000e+00 s per 1e+04

```

```

=====
device name      : TITAN V
device           : cuda:0
Data Management  : Generated on device
Load Time        : 2.365112e-04 s per 1e+05
Transfer Time     : 0.000000e+00 s per 1e+05
Calculation Time  : 1.928306e-02 s per 1e+05
Total Time       : 1.951957e-02 s per 1e+05
Shooby's CPU Time : 9.000000e+00 s per 1e+04

```

```

=====
device name      : TITAN V
device           : cuda:0
Data Management  : Generated on device
Load Time        : 3.108501e-03 s per 1e+06
Transfer Time     : 0.000000e+00 s per 1e+06
Calculation Time  : 1.292557e+00 s per 1e+06
Total Time       : 1.295666e+00 s per 1e+06
Shooby's CPU Time : 9.000000e+00 s per 1e+04

```

In [ ]:

## 1. GPU Time Records (Without Data Transfer)

```

=====
device name      : TITAN V
device           : cuda:0
Data Management  : Generated on device
Load Time        : 1.814365e-03 s per 1e+01
Transfer Time     : 0.000000e+00 s per 1e+01
Calculation Time  : 0.000000e+00 s per 1e+01

```

Total Time : 1.814365e-03 s per 1e+01

```
=====
device name      : TITAN V
device           : cuda:0
Data Management  : Generated on device
Load Time        : 4.076958e-05 s per 1e+02
Transfer Time     : 0.000000e+00 s per 1e+02
Calculation Time  : 0.000000e+00 s per 1e+02
Total Time       : 4.076958e-05 s per 1e+02
=====
```

```
=====
device name      : TITAN V
device           : cuda:0
Data Management  : Generated on device
Load Time        : 5.221367e-05 s per 1e+03
Transfer Time     : 0.000000e+00 s per 1e+03
Calculation Time  : 4.060268e-04 s per 1e+03
Total Time       : 4.582405e-04 s per 1e+03
=====
```

```
=====
device name      : TITAN V
device           : cuda:0
Data Management  : Generated on device
Load Time        : 5.388260e-05 s per 1e+04
Transfer Time     : 0.000000e+00 s per 1e+04
Calculation Time  : 1.125097e-03 s per 1e+04
Total Time       : 1.178980e-03 s per 1e+04
=====
```

```
=====
device name      : TITAN V
device           : cuda:0
Data Management  : Generated on device
Load Time        : 1.611710e-04 s per 1e+05
Transfer Time     : 0.000000e+00 s per 1e+05
Calculation Time  : 1.376677e-02 s per 1e+05
Total Time       : 1.392794e-02 s per 1e+05
=====
```

```
=====
device name      : TITAN V
device           : cuda:0
Data Management  : Generated on device
Load Time        : 3.225565e-03 s per 1e+06
Transfer Time     : 0.000000e+00 s per 1e+06
Calculation Time  : 1.292620e+00 s per 1e+06
Total Time       : 1.295846e+00 s per 1e+06
=====
```

```
In [4]: run1 = 3.5762786865234375e-06, 4.291534423828125e-06, 0.00026679039001464844, 0.001559
run2 = 2.86102294921875e-06, 2.6226043701171875e-06, 0.00026988983154296875, 0.0016419
run3 = 3.0994415283203125e-06, 3.5762786865234375e-06, 0.00048613548278808594, 0.00197
Ns     = np.array([10, 100, 1000, 10000, 100000, 1000000, 10000000])
```

```

run1 = np.array(run1)
run2 = np.array(run2)
run3 = np.array(run3)

run123I = np.mean((run1, run2, run3), 0)

```

## 2. GPU Time Records (With Data Transfer)

```

=====
device name      : TITAN V
device           : cuda:0
Data Management  : Transferred to device
Load Time        : 3.957748e-04 s per 1e+01
Transfer Time     : 6.866455e-05 s per 1e+01
Calculation Time  : 4.634857e-04 s per 1e+01
Total           Time : 9.279251e-04 s per 1e+01
=====

```

```

=====
device name      : TITAN V
device           : cuda:0
Data Management  : Transferred to device
Load Time        : 7.259846e-04 s per 1e+02
Transfer Time     : 4.839897e-05 s per 1e+02
Calculation Time  : 1.916885e-04 s per 1e+02
Total           Time : 9.660721e-04 s per 1e+02
=====

```

```

=====
device name      : TITAN V
device           : cuda:0
Data Management  : Transferred to device
Load Time        : 5.710125e-03 s per 1e+03
Transfer Time     : 2.176762e-04 s per 1e+03
Calculation Time  : 4.482269e-04 s per 1e+03
Total           Time : 6.376028e-03 s per 1e+03
=====

```

```

=====
device name      : TITAN V
device           : cuda:0
Data Management  : Transferred to device
Load Time        : 4.211116e-02 s per 1e+04
Transfer Time     : 1.444340e-03 s per 1e+04
Calculation Time  : 4.163265e-03 s per 1e+04
Total           Time : 4.771876e-02 s per 1e+04
=====

```

```

=====
device name      : TITAN V
device           : cuda:0
Data Management  : Transferred to device
Load Time        : 3.824205e-01 s per 1e+05
Transfer Time     : 1.075053e-02 s per 1e+05
Calculation Time  : 2.390003e-02 s per 1e+05
Total           Time : 4.170711e-01 s per 1e+05
=====

```



```
=====
device name      : TITAN V
device           : cuda:0
Data Management  : Transferred to device
Load Time        : 5.193329e+00 s per 1e+06
Transfer Time     : 1.403193e-01 s per 1e+06
Calculation Time  : 3.657453e-01 s per 1e+06
Total           : 5.699394e+00 s per 1e+06
=====
```

```
In [5]: run1 = 0.00046634674072265625, 0.0010683536529541016, 0.0068149566650390625, 0.05645203
run2 = 0.0008776187896728516, 0.0008144378662109375, 0.006388425827026367, 0.052498102
run3 = 0.0010247230529785156, 0.0011801719665527344, 0.012906789779663086, 0.044039964

run1 = np.array(run1)
run2 = np.array(run2)
run3 = np.array(run3)
```

```
run123II = np.mean((run1, run2, run3), 0)
```

### 3. CPU Time Records

```
=====
device           : cpu Data Management      : Transferred to
device Load Time : 4.427433e-04 s per 1e+01 Transfer
Time            : 5.245209e-06 s per 1e+01 Calculation Time : 1.560926e-03
s per 1e+01 Total Time : 2.008915e-03 s per 1e+01
=====
```

```
device           : cpu Data Management      : Transferred to
device Load Time : 8.444786e-04 s per 1e+02 Transfer
Time            : 9.059906e-06 s per 1e+02 Calculation Time : 1.104832e-02
s per 1e+02 Total Time : 1.190186e-02 s per 1e+02
=====
```

```
device           : cpu Data Management      : Transferred to
device Load Time : 6.634235e-03 s per 1e+03 Transfer
Time            : 2.670288e-05 s per 1e+03 Calculation Time : 9.969068e-02
s per 1e+03 Total Time : 1.063516e-01 s per 1e+03
=====
```

```
device           : cpu Data Management      : Transferred to
device Load Time : 6.788874e-02 s per 1e+04 Transfer
Time            : 2.219677e-04 s per 1e+04 Calculation Time : 1.463918e+00
s per 1e+04 Total Time : 1.532029e+00 s per 1e+04
=====
```

```
device           : cpu Data Management      : Transferred to
device Load Time : 5.639410e-01 s per 1e+05 Transfer
Time            : 1.487732e-03 s per 1e+05 Calculation Time : 1.349072e+01
s per 1e+05 Total Time : 1.405615e+01 s per 1e+05
=====
```

```
device           : cpu Data Management      : Transferred to
device Load Time : 5.423350e+00 s per 1e+06 Transfer
```

```
Time      : 1.540160e-02 s per 1e+06 Calculation Time : 1.374256e+02
s per 1e+06 Total      Time : 1.428644e+02 s per 1e+06
```

```
=====
In [6]: run1 = 0.0007658004760742188, 0.006584882736206055, 0.11615705490112305, 1.33073329925
run2 = run1
run3 = run1
Ns    = np.array([10, 100, 1000, 10000, 100000, 1000000, 10000000])

run1 = np.array(run1)
run2 = np.array(run2)
run3 = np.array(run3)

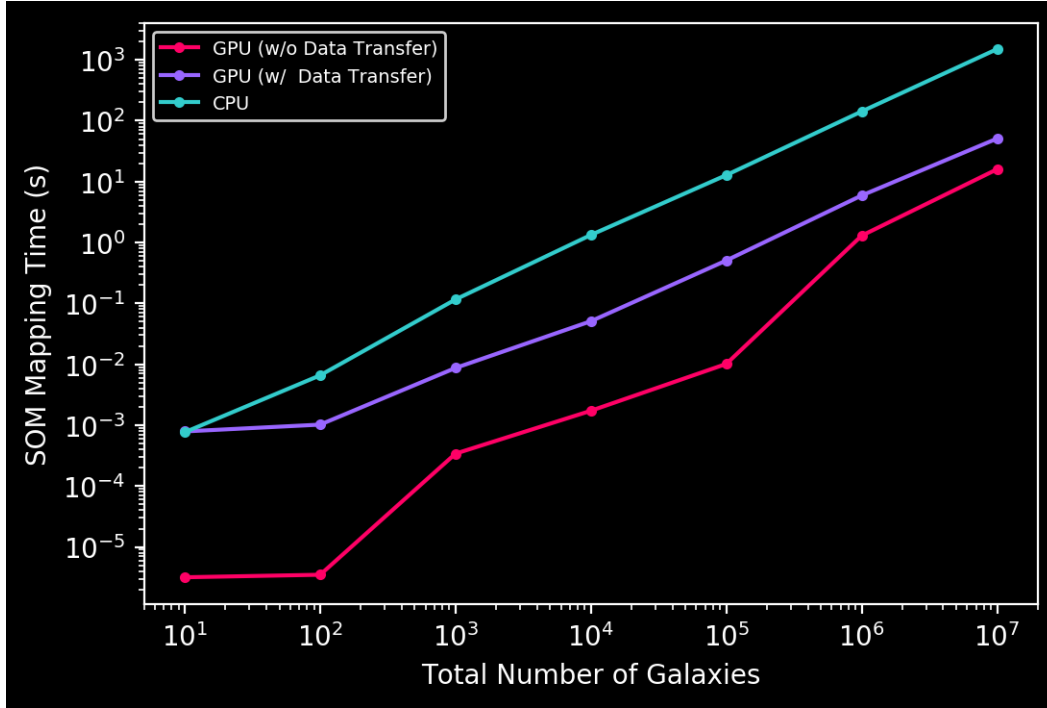
run123III = np.mean((run1, run2, run3), 0)
```

### 0.3.1 Plotting

```
In [11]: import matplotlib.pyplot as plt
         %matplotlib inline
         if True:
             _ = plt.style.use('dark_background')
         fig, axs = plt.subplots(1, 1)
         ax11 = axs

         _ = ax11.loglog(Ns, run123I, label = "GPU (w/o Data Transfer)", marker = '.', color =
         _ = ax11.loglog(Ns, run123II, label = "GPU (w/ Data Transfer)", marker = '.', color =
         _ = ax11.loglog(Ns, run123III, label = "CPU", marker = '.', color = '#33cccc')

         _ = fig.set_dpi(200)
         _ = ax11.set_xlabel('Total Number of Galaxies')
         _ = ax11.set_ylabel('SOM Mapping Time (s)')
         legend = ax11.legend(loc='upper left', shadow=True, fontsize='x-small')
```



## 0.4 Proof of Concept

This section was developed before creating a PyTorch model, even though it is placed last.

### 0.4.1 Data Initialization

```
In [7]: quickTestDict = {'Number of Samples' : 2, 'Vector Size' : 3, 'Grid Height' : 2, 'Grid Width' : 2}
        longTestDict  = {'Number of Samples' : 10, 'Vector Size' : 11, 'Grid Height' : 100, 'Grid Width' : 100}
        paramDict = [quickTestDict, longTestDict][0]

        N          = paramDict['Number of Samples']
        W_shape    = (paramDict['Vector Size'], paramDict['Grid Height'], paramDict['Grid Width'])
        G_shape    = (paramDict['Vector Size'], 1, 1)
        dtype      = np.float32
        rs         = np.random.RandomState(seed=42)
        W           = rs.randint(0, 10, W_shape).astype(dtype)
        Gs          = rs.randint(0, 10, (N, *G_shape)).astype(dtype)

        # Convert to torch tensor
        W           = torch.Tensor(W)
        Gs          = torch.Tensor(Gs)
        print('W Shape: {}'.format(W.shape))
        print('Gs Shape: {}'.format(Gs.shape))
```

```
W Shape: torch.Size([3, 2, 2])
Gs Shape: torch.Size([2, 3, 1, 1])
```

#### 0.4.2 Method I

```
In [8]: if paramDict['dictName'] == 'quickTest':
        for i, G in enumerate(Gs):
            G_grid = torch.ones_like(W)
            G_grid = torch.mul(G, G_grid)
            chi_grid = torch.sum((W - G_grid)**2, 0)
            chi_min = torch.min(chi_grid)
            print("Least Chi-Squared for G[{}]: {:.1e}".format(i, chi_min))
```

```
Least Chi-Squared for G[0]: 4.0e+00
Least Chi-Squared for G[1]: 2.5e+01
```

#### 0.4.3 Method II (Using PyTorch Model)

```
In [11]: print("Model output (min chi for each batch and the corresponding location):")
        model = Model(feed=W)
        model(Gs.to(device))
```

```
Model output (min chi for each batch and the corresponding location):
```

```
Out[11]: (tensor([ 4., 25.], device='cuda:0'), tensor([2, 3], device='cuda:0'))
```

#### 0.4.4 Method III

Compare the output of the following with the Method II output.

```
In [12]: if paramDict['dictName'] == 'quickTest':
        print('\nW = ')
        print(W)
        print('\nGs = ')
        print(Gs)

        print("\n(W - G)**2 for each G:")
        output = torch.add(W, -Gs)**2
        output

        print("\nSum of (W - G)**2 for each G:")
        output.sum(1)
```

```
W =
tensor([[[6., 3.],
```

```

        [7., 4.]],

        [[6., 9.],
         [2., 6.]],

        [[7., 4.],
         [3., 7.]]])

Gs =
tensor([[[[7.]],

        [[2.]],

        [[5.]]],

        [[[4.]],

        [[1.]],

        [[7.]]]])

(W - G)**2 for each G:

Out[12]: tensor([[[[ 1., 16.],
                    [ 0.,  9.]],

                  [[16., 49.],
                    [ 0., 16.]],

                  [[ 4.,  1.],
                    [ 4.,  4.]]],

                [[[ 4.,  1.],
                    [ 9.,  0.]],

                  [[25., 64.],
                    [ 1., 25.]],

                  [[ 0.,  9.],
                    [16.,  0.]])])

```

Sum of  $(W - G)^2$  for each  $G$ :

```

Out[12]: tensor([[[21., 66.],

```

```

[ 4., 29.]],

[[29., 74.],
 [26., 25.]])

```

```

In [13]: # add works on batches too
         if False:
             torch.add(W, -Gs)
             torch.cat([torch.add(W, -Gs[0]), torch.add(W, -Gs[1])])

```

## 0.5 Recycle

```

In [ ]: # Recycled code from torch Model class

# self.G_grid = torch.ones_like(self.W)

#     def ChiFunc1(self, G):
#         G_grid = self.G_grid
#         G_grid = torch.mul(G, self.G_grid)
#         Chi_grid = torch.sum((self.W - G_grid)**2, 0)
#         Chi_min = torch.min(Chi_grid)
#         return Chi_min

#     def ChiFunc2(self, G):
#         output = torch.add(self.W, -G)**2
#         output = torch.sum(output, 0)
#         output = torch.min(output)
#         return output

# Recycled code from Proof of Work

# if True:
#     print('\nW = ')
#     print(W)
#     print('\nG = ')
#     print(G)
#     print('\nW[:, 0, 0] = ', W[:, 0, 0].numpy())
#     print('\nG[:, 0, 0] = ', G[:, 0, 0].numpy())
#     print('\n(G00 - W00)**2 = ', (W[:, 0, 0] - G[:, 0, 0])**2)
#     print('\nsum00 = ', torch.sum((W[:, 0, 0] - G[:, 0, 0])**2))
#     print('\n(G00 - W01)**2 = ', (W[:, 0, 1] - G[:, 0, 0])**2)
#     print('\nsum01 = ', torch.sum((W[:, 0, 1] - G[:, 0, 0])**2))
#     print('\n(G00 - W10)**2 = ', (W[:, 1, 0] - G[:, 0, 0])**2)
#     print('\nsum10 = ', torch.sum((W[:, 1, 0] - G[:, 0, 0])**2))
#     print('\n(G00 - W11)**2 = ', (W[:, 1, 1] - G[:, 0, 0])**2)
#     print('\nsum11 = ', torch.sum((W[:, 1, 1] - G[:, 0, 0])**2))

```